

INTRODUCTION

DMC (Dynamic Models of Choice) is a collection of R functions and associated tutorials written by Andrew Heathcote with contributions from my colleagues (growing out of DE-MCMC code originally written by Brandon Turner and Scott Brown, with stop-signal material contributed by Dora Matzke), postdoctoral fellows (Yishin Lin and Luke Strickland) and students (Quentin Gronau, Angus Reynolds and Matthew Gretton). Its aim is to allow researchers to fit conventional dynamic models of choice (often called “evidence-accumulation” models) using Bayesian methods, and, more importantly, to develop their own new DMCs.

DMCs are hard to work with because they are computationally expensive and have strong correlations among their parameters (they are of the charmingly named class of “sloppy” models common in biology). The aim of DMC is to ease some of those difficulties by providing functions that are convenient and as effective as possible given the challenges, as well as to distill our experience in working with these models into a set of tutorials that provide advice as well as showing how to use the functions.

The DMC tutorials were developed as a means of getting supported hands-on experience in workshops with introductory lectures and recommended background reading. Working through the DMC tutorials by yourself is NOT a good place to start if you are new to Bayesian methods and cognitive models. If that describes you then we recommend you first work through Michael Lee and EJ Wagnemakers wonderful “Bayesian Cognitive Modelling” (<https://bayesmodels.com/>; that was how I learned, by translating the Matlab in an early draft to R). It deals with much more tractable models where things usually work as expected, preparing you for the rigors of DMCs, which are often much less co-operative! More broadly we also recommend “Computational Modelling of Cognition and Behaviour” by Simon Farrell and Steve Lewandowsky. That said, there are lots of notes and references in the tutorials, so if you are more experienced and willing to experiment then you should be fine to work through them by yourself.

Implementing DMC in R fits with the paramount importance of openness, reproducibility, and being able to quickly modify existing models and (relatively!) easily develop new models. However, that means that speed can be an issue; where that becomes an impediment the open source nature of the code means re-implementing in a faster language is always possible. This also has the advantage of the existing DMC as a benchmark. However, we have found that DMC is often sufficient as long as you have access to larger multicore systems, which usually run Linux. DMC was primarily developed for such systems, and that means it usually also works well on Mac, but Windows can be problematic because, at least at the time of writing, the only multicore solution is the Snowfall package, which is outdated and inefficient (especially in terms of memory management).

THE PHILOSOPHY BEHIND DMC

The best way to think of DMC is to be grateful that you didn’t have to implement this stuff from the ground up yourself! It is offered freely but with no guarantees; there is no financial support for the project and no financial benefit to the developers, who are (in the main) cognitive modellers not software engineers, whose main focus is developing new cognitive models and to a lesser degree using already established models (like the DDM or LBA). That means our code is frequently being

updated and likely far from perfect but remember that you are at liberty to improve it, and probably have more time to do so for particular focused projects. We strongly recommend you check everything, and that is easy to do because DMC is completely open.

We have documented DMC through the tutorials specifically because we want you to have to learn the material they contain before you use the system. Our experience is that providing canned solutions in this area just leads to bad science in the form of just accepting estimates or model selection results without checking whether they are valid. To counter this, the idea of “parameter recovery studies”, is built into the bones of DMC. You should always be sure that you can simulate data from your experimental design with your model and then recover the parameters that generated that model before you trust the results of any fits to real data. For some purposes you might have a lower standard “data recovery” (being able to fit the simulated data well but perhaps with some difficulty in getting the parameters back), but even in this case the point is you need to simulate and fit in order to know where you stand. We have a paper that reflects this philosophy and provide examples that address more advanced uses of DMC (at the time of writing still in press). If you have time it is worth a read first to get an overview, but the main thing to do is work through the tutorials

The philosophy behind DMC is also reflected in how you add new models, which requires specifying a random function (that simulates data), not just a likelihood function (that allows you to fit data), and in the tutorials, which are all about simulating data and then fitting it. Indeed it is not until the end of the first part of lesson 4 (`dmc_4_1_Simulating.R``) that we explicitly tell you how to use DMC with real data! We strongly recommend that you finish to the end of lesson 4 before you think about fitting your own data. At that time, you might also want to work through the advanced lessons that are flagged throughout the first four (mainly those in the lesson 5 series). The first four lessons address three standard models; for those interested in going beyond that, a few new models are developed in the lesson 6 series. If you want to add your own models to DMC then read lessons 2.3 and 2.4 first.

USING DMC

=====

Before using DMC go into the packages directory, open ``install_packages.R`` and make sure all of the standard CRAN packages are available, as well as installing the tweaked version of the coda package. Many lessons have associated data files that contain the results that are summarised in the lesson. When you run commands you will get different results, so for comparison (or just a first run through because it can often take quite a while to run the commands) you can load these data files: ``load_data`` commands to do so are in the lessons but commented out. Note that if the RData objects are not in the tutorial/data subdirectory ``load_data`` will fetch them from the internet. If you want to have all data files local from the start use the ``get_all_data.R`` script in the tutorial subdirectory. Note also that while we try to keep the comments in the tutorials congruent with the stored data files that sometimes slips, sorry ...

We welcome bug reports (via email to myself) and like to help users, but our capacity to do so is limited. We are constantly updating DMC, mainly with the aim of improving its ability to implement new models, and that can sometimes cause instabilities for old models. Old versions of DMC are archived on the OSF site (<https://osf.io/pbwx8/>; available under Release History in the Files

section), so take note of the version you use with each project. OSF also provides a facility to archive a self-contained version of your work (i.e., the particular version of the base DMC code you used, your modifications, and any scripts, parameters recovery studies, fits and associated data etc.), so that others can verify it or modify it to use in their own projects. This can be achieved via forking the DMC project to your own OSF account (via the 'fork' button at the top right of the project page).

The model directories included in the current version are (see also "TutorialModels.txt" in the dmc/models directory, others are being added constantly, sometimes we slip and include them in a release):

DDM, EXG-SS, LBA, LBA-GoFailure, LBA-GoNoGo, LBA-pda, LBA-T0Dists,
LNR, LNR-GoNoGo, LNR-pda, Wald

Andrew Heathcote (andrew.heathcote@utas.edu.au)

Amsterdam, 11/July/2019

Tables of DMC lessons and their contents.

| Tutorial | Topic | New functions |
|-----------------|---|---|
| dmc_1_1 | Set up a DMC model object | load_model Loads a model class into the R environment. model.dmc Creates a dmc model object defining the design of the data. print.cell.p Prints parameter values for each cell of a design. |
| dmc_1_2 | Simulating and exploring the LNR/LBA/DDM. | load_data Loads data stored in the data folder. simulate.dmc Simulates fake data from a model and parameter vector. data.model.dmc Converts data into a data-model object that stores additional information about the model necessary for Bayesian sampling. likelihood.dmc Returns the likelihood of the data given a parameter vector. plot.cell.density Plot the response time density of each response for a given cell of the design. profile.dmc Create profile plots (likelihood curves) of data by varying one parameter in the parameter vector. |
| dmc_1_3 | Build an LBA model | Introduces use of the lba_B.R model file in models/LBA. plot.score.dmc Returns the response accuracy and the mean RT of correct and error responses. |
| dmc_1_4 | Build an LNR model | Introduces use of the lnr.R model file in models/LNR. |
| dmc_1_5 | Build a DDM | Introduces use of the ddm.R model file in models/DDM. |

| Tutorial | Topic | New functions |
|----------|-------------------------------------|--|
| dmc_2_1 | Prior distributions - Basic | <p>prior.p.dmc Generate a prior object.</p> <p>plot.prior Plot the prior distributions.</p> <p>check.recovery.dmc Returns a parameter recovery summary, including the difference between the true value and the median sampled value if the true values are given.</p> <p>summary.dmc More detailed information about parameter estimates.</p> |
| dmc_2_2 | Prior distributions - Advanced | <p>log.prior.dmc Calculates the log of the value of the prior density of a set of parameter values.</p> <p>log.posterior.dmc Calculates the summed log posterior likelihood, which is the sum of the log likelihood and the sum of the log prior given specific values of the parameters.</p> |
| dmc_2_3 | Adding new models to DMC | <p>random.dmc Internal functions used by <i>simulate.dmc</i> to generate random data.</p> <p>transform.dmc Internal functions used to transform parameter values to match with the likelihood function. For example, B vs b notation in the LBA.</p> |
| dmc_2_4 | Adding new models to DMC - Advanced | <p>rlba.norm Generate LBA data with a normal distributed rate. Based off the rtdists package <i>rlba_norm</i>.</p> <p>p.df.dmc Internal function that manages parameter values.</p> |

| Tutorial | Topic | New functions |
|-----------------|-----------------------------------|---|
| dmc_3_1 | Fit an LNR model | samples.dmc Creates an object used to control the sampling process. run.dmc Runs sampling iterations. plot.dmc Plots DE-MCMC chains. If given a prior object it will compare the prior distributions to the posterior distributions. rprior.dmc Sample from a constant prior. acf.dmc Plot the auto-correlation of subsequent samples. |
| dmc_3_2 | Assessing the fit of an LNR model | pick.stuck.dmc Identify stuck chains. gelman.diag.dmc Returns Gelman's diagnostic values, used to determine if the samples have converged. effectiveSize.dmc Return the effective sample size of the samples (taking into account the auto-correlations). post.predict.dmc Generate predicted data from a posterior. plot.pp.dmc Plot density of posterior prediction data and compare to original data. pairs.dmc Plot scatters of each possible pair of parameters after sampling. Used to identify correlated parameters. run.unstuck.dmc Automatically repeat sampling until there are no stuck chains. run.converge.dmc Automatically run sampling until some convergence criteria are met. |
| dmc_3_3 | Fit an LBA model | |
| dmc_3_4 | Fit a DDM model | ppp.dmc Posterior predictive p-values. |

| Tutorial | Topic | New functions |
|----------|------------------------------------|--|
| dmc_3_5 | Individual subject model selection | <p>Dstats.dmc Returns a list of statistics about the posterior deviance.</p> <p>pd.dmc Estimates the “true” number of parameters.</p> <p>plot.deviance.dmc Plots posterior deviance and shows estimates of the true number of parameters.</p> <p>posterior.lr.dmc Performs a posterior likelihood ratio test.</p> <p>IC.dmc Calculates DIC in different ways.</p> <p>wIC.dmc IC tests with different weights.</p> <p>trial_log_likes Returns log likelihoods of samples with thinning.</p> <p>waic.dmc Computes Watanabe-Akaike information criterion.</p> <p>loocompare.dmc Makes comparisons on two or more WAIC results.</p> <p>looic.dmc Computes LOO Information criterion.</p> <p>p.fun.dmc Do posterior predictive tests on functions of parameters (e.g., the difference between two parameters).</p> |

| Tutorial | Topic | New functions |
|-----------------|--|--|
| dmc_4_1 | Conduct model simulation for multiple subjects | h.simulate.dmc Simulate hierarchical data. |
| dmc_4_2 | Prior distributions – Hierarchical models | |
| dmc_4_3 | Posterior distributions – Hierarchical models | |
| dmc_4_4 | Fit a fixed effect LNR model | h.log.likelihood.dmc Calculate the log-likelihood of data given a hierarchical model. assign.pp Convenience function used for hierarchical profile plots. h.profile.dmc Graph hierarchical profile plots. h.samples.dmc Set up for hierarchical sampling. h.run.dmc Run hierarchical sampling iterations. h.pick.stuck.dmc Pick out stuck chains. h.run.unstuck.dmc Automatically run hierarchical sampling until there are no stuck chains. h.run.converge.dmc Automatically run hierarchical sampling until it reaches convergence criteria. group_trial_log_likes Calculates and stores likelihoods at the trial level. group_subject_log_likes Calculates and stores likelihoods at the subject level. h.post.predict.dmc Generate posterior predictive data from a hierarchical model. h.check.recovery.dmc Return information about parameter recovery. |
| dmc_4_5 | Fit a random effect LNR model | h.gelman.diag.dmc Calculate Gelman’s diagnostic on hierarchical samples. |

| Tutorial | Topic | New functions |
|----------|---|--|
| dmc_4_6 | Fit fixed- and random-effect LBA models | get.thin Estimates how much thinning to do. make.hstart Makes hierarchical start points from individual fits make.theta1 Extracts last sample from individual fits. |
| dmc_4_7 | Fit fixed- and random-effect DDM models | |
| dmc_5_1 | Complex factorial designs | check.p.vector Make sure the parameter vector matches the design of the model. empty.map Makes a map object used to define non-standard design. assign.map Assigns cells to map object. h.IC.dmc Hierarchical version of <i>IC.dmc</i> . cor.plausible Calculates plausible correlation values. postRav Makes summary of plausible values over subjects. postRav.Density Extract density from <i>postRav</i> output. postRav.mean Extract mean from <i>postRav</i> output. postRav.p Extract probability < 0 from <i>postRav</i> output. postRav.ci Extract credible interval from <i>postRav</i> output. compare.r Compares two sets of plausible value correlations. |
| dmc_5_2 | Advanced scoring of accuracy | |
| dmc_5_3 | Hierarchical model selection | |
| dmc_5_4 | Plausible values | |

| Tutorial | Topic | New functions |
|-----------------|--|--|
| dmc_5_5 | Advanced plotting | ggplot.RP.dmc Plots the proportion of each response. ggplot.RT.dmc Plots RT quantiles. get.fitgglist.dmc Gets a data frame of summary stats for ggplot. ggplot.RA.dmc Plots accuracy. |
| dmc_5_6 | Testing parameter effects | compare.p Within-subjects posterior inference. compare.ps Between-subjects posterior inference. |
| dmc_5_7 | Bayes Factors | bridge.sampler.dmc Calculates marginal likelihoods from samples object bf.dmc Calculates Bayes Factors from marginal likelihoods |
| dmc_6_1 | Fit LBA model with go failure | Introduces use of lba_Bgf.R model file in models/LBA-GoFailure. |
| dmc_6_2 | Fit Wald models with go failure | Introduces use of wald.R model file in models/Wald. |
| dmc_6_3 | Fit Go-NoGo models | Introduces use of lnrgng.R model file in models/LNR-GoNoGo and lba_Bgng.R in models/LBA. |
| dmc_6_4 | Fit the ExGaussian stop-signal model with two racers | Introduces use of exgSS.R model file in models/EXG-SS. plot_SS_go.dmc Plot go RTs. plot_SS_if.dmc Plot inhibition function. plot_SS_srrt.dmc Plot signal-respond RTs. |
| dmc_6_5 | Fit the ExGaussian stop-signal model with three racers | Introduces use of exgSSprobit.R model file in models/EXG-SS. |
| DMCpaper1 | Advanced fitting and assessment of LBA models | RUN.dmc Repeatedly run <i>run.converge</i> and <i>run.unstuck</i> until all the convergence criteria are met. h.RUN.dmc Same as above but with multiple subjects. |
| DMCpaper2 | Advanced fitting and assessing of stop-signal models | |

